

PID controllers in DeFi: technical architecture, vulnerabilities, and implementation

PID-controlled stablecoins like RAI represent one of DeFi's most sophisticated control theory applications, using feedback mechanisms to maintain price stability without hard pegs. [Medium ↗](#) RAI's Proportional-Integral (PI) controller—deliberately omitting the derivative term—has demonstrated remarkable effectiveness: [Medium ↗](#) during ETH's 350% price increase from October 2020 to January 2021, RAI's redemption price fluctuated less than 4%. [DeFi Pulse ↗](#) This stability comes from careful parameter tuning via cadCAD simulations, anti-windup mechanisms using leaky integrators, and oracle manipulation resistance through 12-hour TWAP windows. However, implementing PID controllers on-chain introduces unique challenges around gas costs, fixed-point arithmetic precision, and novel attack vectors that don't exist in traditional control systems.

RAI's PI controller architecture sets the standard

Reflexer Labs' RAI implements a **PI controller** (not full PID) that computes a redemption rate based on the deviation between market price and an internal redemption price. The mathematical foundation is straightforward:



$$\text{redemption_rate} = K_p \times (\text{market_price} - \text{redemption_price}) + K_i \times \int(\text{error})dt$$

The derivative term (K_d) was intentionally set to zero because market prices contain significant noise that derivative calculations would amplify, causing erratic controller behavior. This design decision, validated through extensive BlockScience simulations, demonstrates that **simpler control architectures often outperform theoretically optimal ones** in adversarial blockchain environments.

The controller's core feedback loop operates as follows: when market price exceeds redemption price, the system sets a negative redemption rate, causing RAI's internal valuation to depreciate. [Dankrad Feist ↗](#) This creates arbitrage incentives—users can mint RAI at the lower redemption price and sell at the higher market price, creating selling pressure that pushes prices back toward equilibrium. [TokenBrice ↗](#) [The Defiant ↗](#) The inverse occurs when market price falls below redemption price. [Medium ↗](#)

RAI's **K_p term operates 3+ orders of magnitude stronger than K_i**, ensuring proportional response dominates while the integral term slowly eliminates steady-state errors. The integral uses a "leaky integrator" that decays approximately **1/1000 of its value per period**, [Medium ↗](#) preventing integral windup where accumulated error could overpower proportional control during sustained price deviations.

Smart contract implementation spans multiple modules

RAI's PID logic is distributed across several contracts in the GEB (Generalized Ethereum Bindings) framework:

Contract	Function
PIRawPerSecondCalculator	Core PI calculation using raw price deviation
PIScaledPerSecondCalculator	PI using scaled deviation: `
PIRateSetter	Orchestrates calculator calls, passes results to relayer
OracleRelayer	Stores redemption price/rate, interfaces with oracles
UniswapConsecutiveSlotsPriceFeedMedianizer	12-hour TWAP oracle for market price

The OracleRelayer continuously updates redemption price using exponential accumulation:



```
redemptionPrice = rmultiplication(
    rpower(redemptionRate, block.timestamp - redemptionPriceUpdateTime, RAY),
    _redemptionPrice
);
```

This calculation uses **27-decimal (RAY) precision** for rates while most token amounts use 18-decimal (WAD) precision. The rpower function implements repeated squaring for gas-efficient exponentiation, handling arbitrary time deltas between updates.

HAI and alternative implementations expand the design space

HAI, deployed on **Optimism in February 2024**, forks RAI's PI controller architecture while introducing multi-collateral support (ETH, wstETH, OP, WETH, tBTC, VELO). This addresses RAI's primary limitation: ETH-only collateral creates opportunity costs for users wanting staking yields.

Key technical differences between HAI and RAI:

Aspect	RAI	HAI
Initial Price	\$3.14 (symbolic)	\$1.00 (practical)
Collateral	ETH only	6+ collateral types
Chain	Ethereum L1	Optimism L2
Governance	Minimized (FLX)	Active (KITE)
Gas Costs	High (~\$5-50 per operation)	Low (<\$0.10)

HAI's L2 deployment enables **more frequent controller updates** and lower minimum vault sizes, since gas costs don't prohibit small operations. The multi-collateral design with yield-bearing assets like wstETH is expected to produce **positive equilibrium redemption rates** (HAI appreciates over time), contrasting with RAI's historically negative rates.

Beyond stablecoins, **Euler Finance** applies PID controllers to interest rate markets. Their Adaptive Curve IRM automatically adjusts borrowing rates based on utilization deviation from targets, amplifying or dampening rate changes when utilization strays from optimal levels. Research has shown these PI-based interest controllers are vulnerable to **just-in-time liquidity attacks** where sophisticated actors manipulate utilization to extract excess yield.

Vulnerability landscape spans oracles, governance, and arithmetic

PID controllers in DeFi introduce attack surfaces that don't exist in traditional control systems. While RAI has not suffered major exploits, OpenZeppelin's November 2020 audit uncovered several high and medium severity issues.

Oracle manipulation poses the greatest risk

The controller's accuracy depends entirely on price feed integrity. A **1% deviation** between TWAP price and redemption price creates approximately **±6% annual redemption rate impact**, meaning oracle manipulation has leveraged effects on system behavior.

Attack vectors include:

- **TWAP manipulation:** Acquiring significant position, trading to move price, holding position through TWAP window
- **Stale price exploitation:** Acting on information before oracle updates propagate
- **Flash crash returns:** Extreme values during brief price dislocations triggering aggressive rate changes

RAI mitigates these through minimum liquidity requirements (**\$2M on oracle exchanges**, at least **3% of supply** on monitored exchanges) [Reflexer ↗](#) and extended TWAP windows (12 hours). However, these mitigations trade security for responsiveness—the 12-hour lag creates arbitrage opportunities for informed traders.

Integral windup enables whale manipulation

BlockScience's "Moby Dick" attack analysis showed that a well-capitalized attacker could buy large RAI quantities and hold market price constant. Without proper anti-windup mechanisms, the integral term accumulates unboundedly, eventually causing explosive rate changes. RAI's leaky integrator provides analytical bounds ensuring the integral cannot overpower proportional control—[Medium ↗](#) the attacker eventually loses money as the proportional response dominates. [Medium ↗](#)

Governance attacks can manipulate parameters

Similar to the **\$182M Beanstalk exploit**, flash-loaned governance tokens could pass malicious proposals [OxJournal ↗](#) modifying Kp/Ki values:

- **Kp too high:** System oscillates, potentially triggering cascading liquidations
- **Kp too low:** Controller ineffective, market price diverges persistently
- **Ki too high:** Integral windup despite anti-windup mechanisms
- **Ki sign inverted:** Positive feedback loop destabilizes system

RAI implements timelocks on parameter changes and bounds on redemption rates at the OracleRelayer level as final safeguards. [Reflexer ↗](#)

Technical vulnerabilities from audit findings

The OpenZeppelin audit identified critical issues including:

- **[H01] Debt accounting error** in `fastTrackAuction` during global settlement, skewing collateral valuations
- **[M01/M04] DoS vectors** allowing front-running of debt/surplus auctions
- **[M06] Unsafe integer casting** causing overflow for large values
- **[M09] WAD rounding errors** in collateral auctions

On-chain computation imposes fundamental constraints

Implementing control algorithms on Ethereum requires navigating significant technical limitations around gas costs, arithmetic precision, and time handling.

Fixed-point arithmetic replaces floating-point

The EVM only supports 256-bit integers, requiring all fractional calculations to use scaled integer arithmetic:

Precision Level	Scale Factor	Use Case
WAD	10^{18}	Token amounts, basic calculations
RAY	10^{27}	Interest rates, redemption rates
UQ112x112	2^{112}	Uniswap V2 price accumulators

Rounding errors accumulate through operations. Best practices include:

- **Multiply before divide** to preserve precision: $(a \times b) / c$ instead of $(a / c) \times b$
- **Directional rounding** toward protocol (against users): `mulWadUp()`, `divWadUp()` variants
- **Intermediate precision** using RAY for calculations, converting to WAD for results

Gas costs constrain update frequency

Storage operations dominate PID update costs:

- **SLOAD (cold)**: 2,100 gas
- **SSTORE (new)**: 20,000+ gas
- **Total PID update**: ~50,000-150,000 gas depending on complexity

At \$50 gwei and \$3,000 ETH, a 100,000 gas PID update costs ~\$15. This makes frequent updates economically infeasible on L1—RAI's 12-hour minimum update interval partially reflects this constraint. HAI's Optimism deployment reduces costs by **90%+**, enabling more responsive control.

Block time variability requires careful handling

Controllers use per-second rates with exponential accumulation to handle arbitrary time gaps:



solidity

```
// Works regardless of time since last update
new_rate = old_rate × (1 + per_second_fee)^seconds_elapsed
```

MakerDAO's `drip()` function pioneered this pattern, and RAI adopts it for redemption price updates. However, debt created and repaid between updates incurs zero stability fees—keeper incentives help ensure regular updates.

Minimal viable implementation requires fewer components than expected

A production PID controller can start surprisingly simple. RAI launched with **P-only control** (no integral term), demonstrating that minimal architectures can work effectively before upgrading.

Essential vs optional components for MVP

Component	P-Only MVP	Production PI
Price oracle (TWAP preferred)	Required	Required
Redemption price storage	Required	Required
Proportional gain (K_p)	Required	Required
Integral accumulator	Not needed	Required
Anti-windup mechanism	Not needed	Critical
Rate bounds	Recommended	Required
Circuit breakers	Recommended	Required
Governance timelock	Recommended	Required

A minimal P-controller implementation requires only ~50 lines of Solidity:



```
contract MinimalPController {  
    uint256 public redemptionPrice = 1e18; // Start at $1  
    uint256 public Kp = 1e15; // 0.001 gain (tunable)  
    uint256 public lastUpdate;  
  
    function updateRate(uint256 marketPrice) external {  
        int256 error = int256(marketPrice) - int256(redemptionPrice);  
        int256 rateAdjustment = (error * int256(Kp)) / 1e18;  
  
        uint256 elapsed = block.timestamp - lastUpdate;  
        redemptionPrice = uint256(  
            int256(redemptionPrice) + (rateAdjustment * int256(elapsed))  
        );  
        lastUpdate = block.timestamp;  
    }  
}
```

Safety mechanisms are non-negotiable

Even minimal implementations require:

- **Rate bounds:** Prevent extreme redemption rate changes

- **Staleness checks:** Reject oracle prices older than threshold (e.g., 1 hour)
- **Pause capability:** Emergency stop for detected manipulation
- **Minimum liquidity requirements:** Don't activate controller without sufficient exchange depth

RAI's documentation explicitly warns: PID should NOT be set to full force at launch. [Reflexer ↗](#) Conservative initial parameters with gradual strengthening as confidence grows is the established pattern.

USDC-based systems enable simplified designs

Applying PID controllers to USDC-collateralized systems offers advantages over volatile collateral like ETH:

Aspect	ETH-Backed	USDC-Backed
Collateralization ratio	135%+ required	110-120% possible
Oracle complexity	ETH/USD volatility	USDC ≈ \$1 (simple)
Liquidation risk	Higher	Lower
Censorship risk	None	USDC blacklisting possible

Interest rate PID controllers for USDC lending markets represent the most practical application. Rather than controlling a stablecoin's price, the PID maintains target utilization rates:



$$\text{Error} = \text{Actual_Utilization} - \text{Target_Utilization}$$

$$\text{Interest_Rate} = \text{Base_Rate} + K_p \times \text{Error} + K_i \times \int (\text{Error}) dt$$

This eliminates governance decisions about optimal interest rates—the controller automatically adjusts to maintain 80% utilization (or any target), improving capital efficiency while reducing governance surface area.

For USDC-backed synthetic assets, lower collateral buffers are possible because USDC volatility is minimal. However, **USDC's blacklisting capability** introduces censorship risk that doesn't exist with ETH collateral. Any system design must handle the scenario where collateral addresses are frozen.

Conclusion

PID controllers represent a maturing technology in DeFi, with RAI's **3+ years of mainnet operation** proving the viability of control-theory-based stability mechanisms. Key insights for implementers:

Start simpler than you think necessary. RAI succeeded with P-only control initially; derivative terms are almost never worth the complexity in blockchain environments where price data is inherently noisy.

Anti-windup is critical for any integral term. The leaky integrator pattern—decaying accumulated error at ~0.1% per period—prevents sophisticated attackers from exploiting integral windup even under sustained manipulation. [Medium ↗](#)

L2 deployment dramatically improves economics. HAI's Optimism deployment enables update frequencies and vault sizes impossible on L1, suggesting future PID implementations should target L2 by default.

Oracle design is the primary security consideration. The controller's accuracy depends entirely on price feed integrity; extended TWAP windows (12+ hours) and minimum liquidity requirements (\$2M+) provide manipulation resistance at the

cost of responsiveness. [Reflexer](#)↗

The technology has proven itself for floating-peg stablecoins and shows promise for interest rate control in lending markets. For teams considering implementation, the path is clear: begin with a minimal P-controller, validate behavior through testnet deployment and simulation, then upgrade to PI control with proper anti-windup once operational stability is established.